

# Database Cracking: Concept Evaluation

Zulsar Batmunkh  
MIT

Dana Mukusheva  
MIT

**Abstract**—abstract goes here.

## 1 INTRODUCTION

## 2 BACKGROUND

### 2.1 Database Cracking Overview

Database cracking is a concept of adaptive column partitioning and sorting based on the query workloads and patterns. The core idea of database cracking is to partition the column according to each incoming query, therefore, effectively reducing the size of partitions as more queries arrive. Such self-organizational behavior is desired when there is no upfront knowledge of query patterns, which is required to preliminary create traditional indexes of the columns of interest.

### 2.2 Previous work

Database cracking is an active research topic, which is led by a group of researchers from CWI Amsterdam. The concept of self-organizing columns and query-dependent column partitioning was first presented in by Kersten and Manegold [2]. Blablabla one Nobody [3]. Blablabla two Nobody [1]. Blablabla three Nobody [2].

### 2.3 Cracking Algorithms

## 3 SYSTEM OVERVIEW

In order to analyze the performance effects of the database cracking, we have implemented our own simple database system MiniDB. In this section we will describe components and implementation details of MiniDB.

MiniDB is a single column database implemented in Java, where each table consists of a single uniquely named column. The database keeps a mapping between a column name and a column object in a hash map. The database table stores tuples as a list, and each tuple is a 32-bit integer. MiniDB maintains the following data structures: Simple Column, Sorted Column, Cracker Column, Cracker Index, Range Scan.

### 3.1 MiniDB columns and indexes

Simple Columns represent database tables and store tuples. They do not maintain a specific tuple order and insert a tuple to the end of the tuple list. In Simple Columns tuple lookup requires linear scan of the tuples list.

Sorted Column is another data structure that represents database tables which preserves the order of the tuples. Tuple lookup takes logarithmic time and is implemented as a binary search.

Cracker Column is a data structure that cannot exist independently in the database, it can only be coupled with Simple Column that supports cracking. They are initialized and attached to a Simple Column instance after the first query. Cracker Columns contain same values as their corresponding Simple Columns, but in a different and constantly changing order. Cracker Columns store tuples in a partially sorted list, that is, once its tuples are reorganized and partitioned into two sublists (one with all tuples whose values are less than or equal to the partition value and one with all tuples whose values are greater than the partition value) Each Cracker Column is supplemented with a Cracker Index instance.

	AVL Tree	Hash Map	Sorted List
<b>Pros</b>	Logarithmic insert, lookup, predecessor and successor search	Constant time lookup	Logarithmic lookup and constant predecessor and successor search
<b>Cons</b>	Cost for balance maintenance	Linear predecessor and successor search	Linear insert

Table 1. Summary of cost-related advantages and disadvantages for each Cracker Index implementation.

The Cracker Index instances are the data structures that are necessary to keep most-up-to date information about all partitions of the Cracker Column tuples. Particularly, Cracker Index stores pairs in the form  $(v, p)$ , where  $v$  indicates that all tuples located at the positions less than and including  $p$  have values less than and including  $v$ .

### 3.2 Query processing

MiniDB queries are the Range Scan objects that operate on a single column. Each Range Scan instance stores the pointer to its column of interest, endpoints of the value ranges, and the range sign, either one of  $\leq, <, \geq, >, <=, <=<, <=>, <=>=<=>$ . Range Scan objects are essentially iterators on the values of their columns of interests. If the column does not support cracking, the iterator goes over all tuples and returns only those whose values belong to the specified range. Otherwise, they use cracking and simply iterate over all values that lie in a specified partition.

### 3.3 Cracker Index implementation

Cracker Index functionality requires methods such as insert value-position pair, find predecessor of value  $v$ , find successor of value  $v$  and lookup position  $p$  of value  $v$ . Predecessor/successor search is used to define partitions with the closest boundaries, which is the case when a new unseen query arrives. We have implemented Cracker Index using three different underlying representations, AVL Tree, Sorted List and HashMap. Each of them has advantages and disadvantages in different scenarios.

**AVL Tree** stores value  $v$  as a node key and a position  $p$  as a node data. All of the operations on AVL Tree have logarithmic cost, which makes it a good candidates for large workloads. However, at the same time, the size of the tree grows with a number of queries, and maintaining the balance of the large tree might be costly.

**Hash Map** directly stores mapping between the value  $v$  and the position  $p$ . The HashMap implementation is beneficial when all incoming queries are repetitive, and additional cracking is unnecessary, since the correct partition already exists. However, adding a new partition info into the index is costly, as predecessor/successor search takes linear time.

**Sorted List** stores value-position pairs a list sorted on values. Preserving the order of the list is costly and takes linear time, however, cheap successor and predecessor search balances the query cost.

- Zulsar Batmunkh is MIT undergraduate. email: zulsar@mit.edu
- Dana Mukusheva is MIT undergraduate. email: mukushev@mit.edu

In Table 1 we summarized pros and cons for each Cracker Index implementation.

## **4 EXPERIMENTS**

### **4.1 Experiment Setup**

Describe the hardware (server, memory, OS, number of cores, etc)  
Describe how we timed things

### **4.2 Performance Evaluation**

### **4.3 Comparison to MonetDB**

not sure if this comes here

## **5 DISCUSSION**

Talk about what needs to be done to improve the quality of our experiments, some assumptions we made and what we have neglected

## **6 CONCLUSION**

## **REFERENCES**

- [1] S. Idreos, M. L. Kersten, and S. Manegold. Database cracking. In *CIDR*, pages 68–78, 2007.
- [2] M. Kersten and S. Manegold. Cracking the database store. In *CIDR*, 2005.
- [3] F. M. Schuhknecht, A. Jindal, and J. Dittrich. The Uncracked Pieces in Database Cracking Management. *Proc. VLDB Endow.*, 7(2):97–108, Oct. 2013.